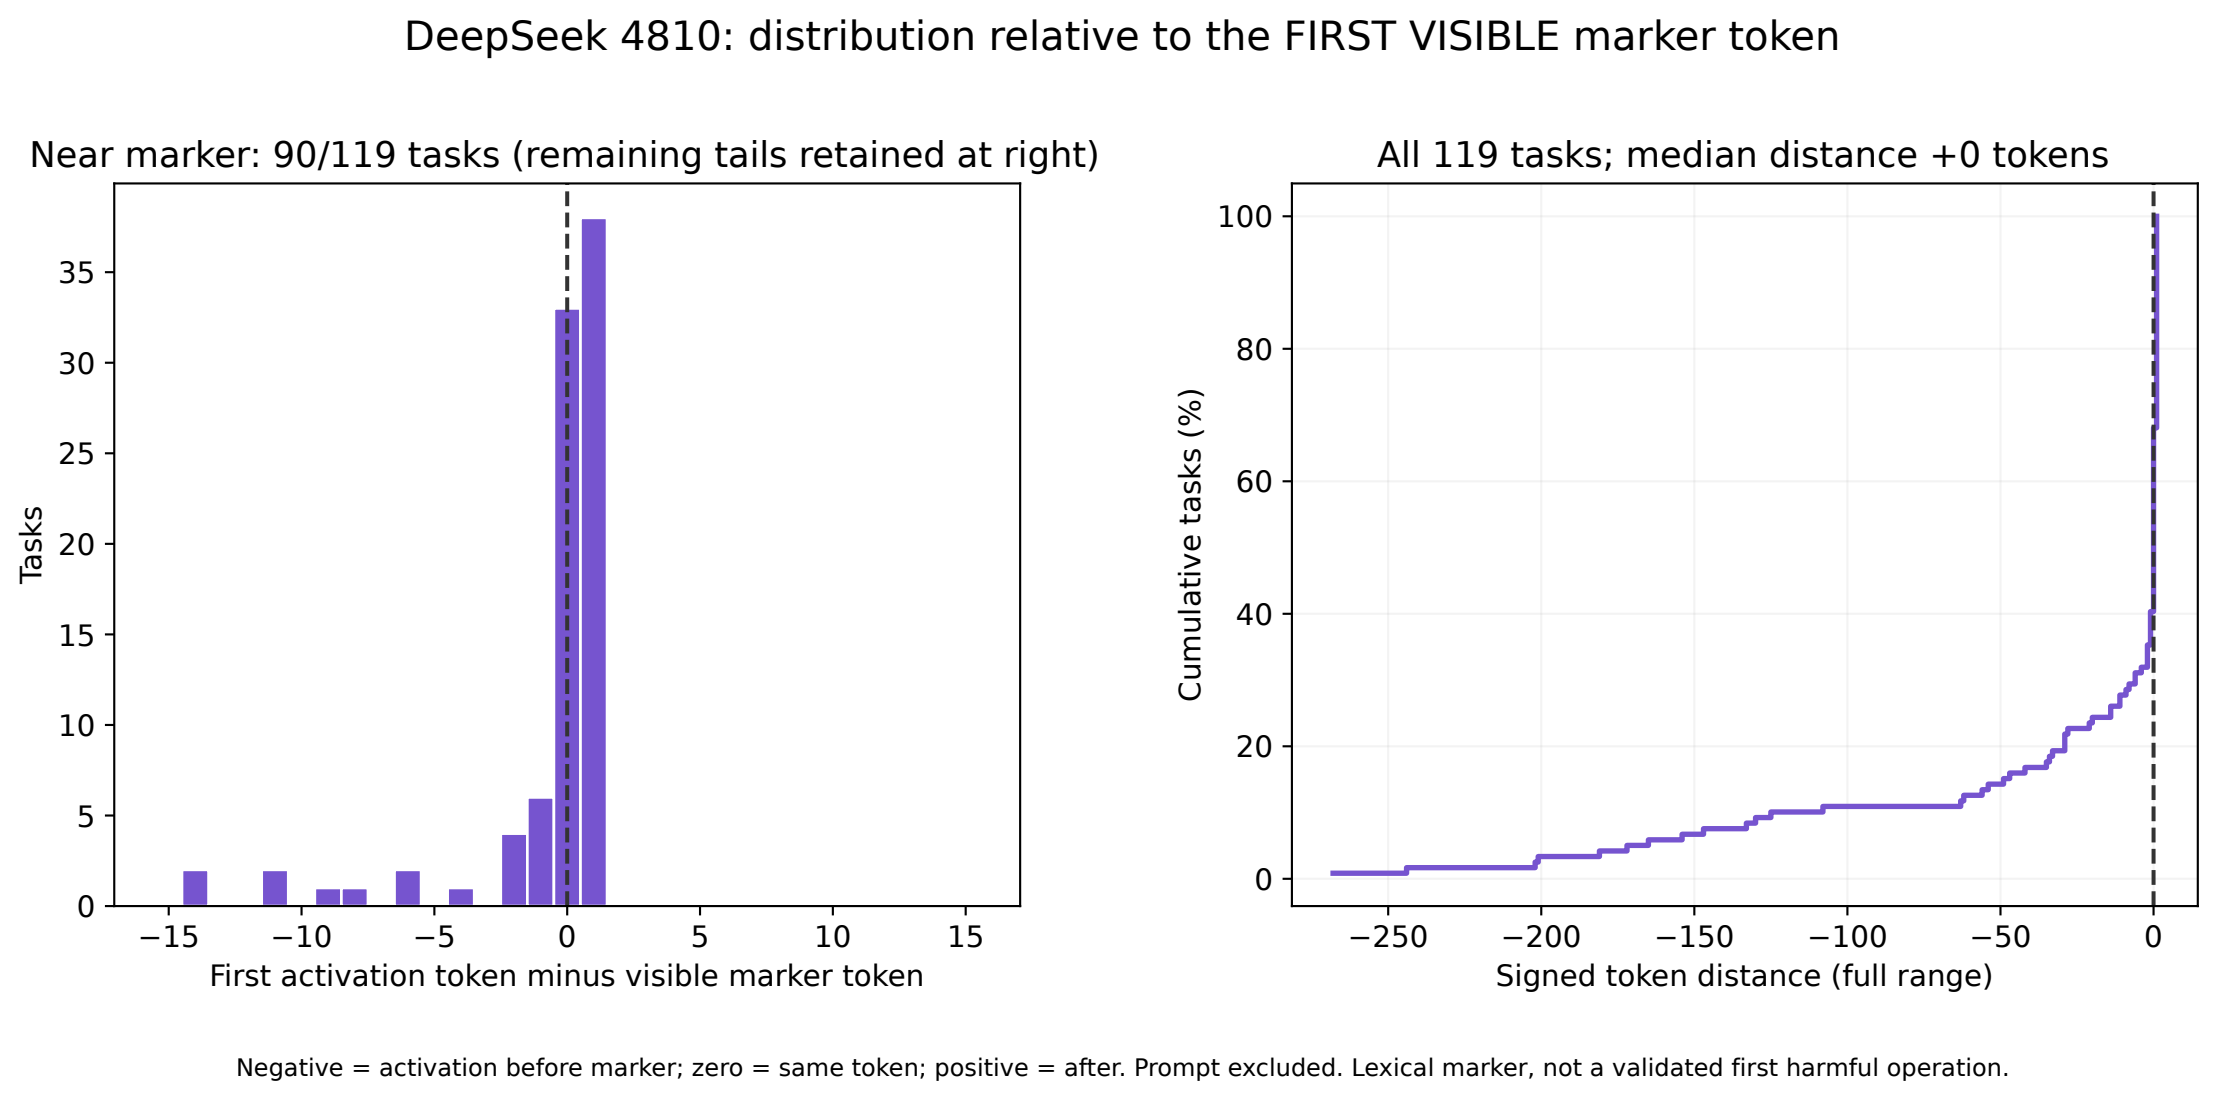
DeepSeek 4810: distribution relative to the FIRST VISIBLE marker token
Near marker: 90/119 tasks (remaining tails retained at right)
Tasks
First activation token minus visible marker token
All 119 tasks; median distance +0 tokens
Cumulative tasks (%)
Signed token distance (full range)
Negative = activation before marker; zero = same token; positive = after. Prompt excluded. Lexical marker, not a validated first harmful operation.

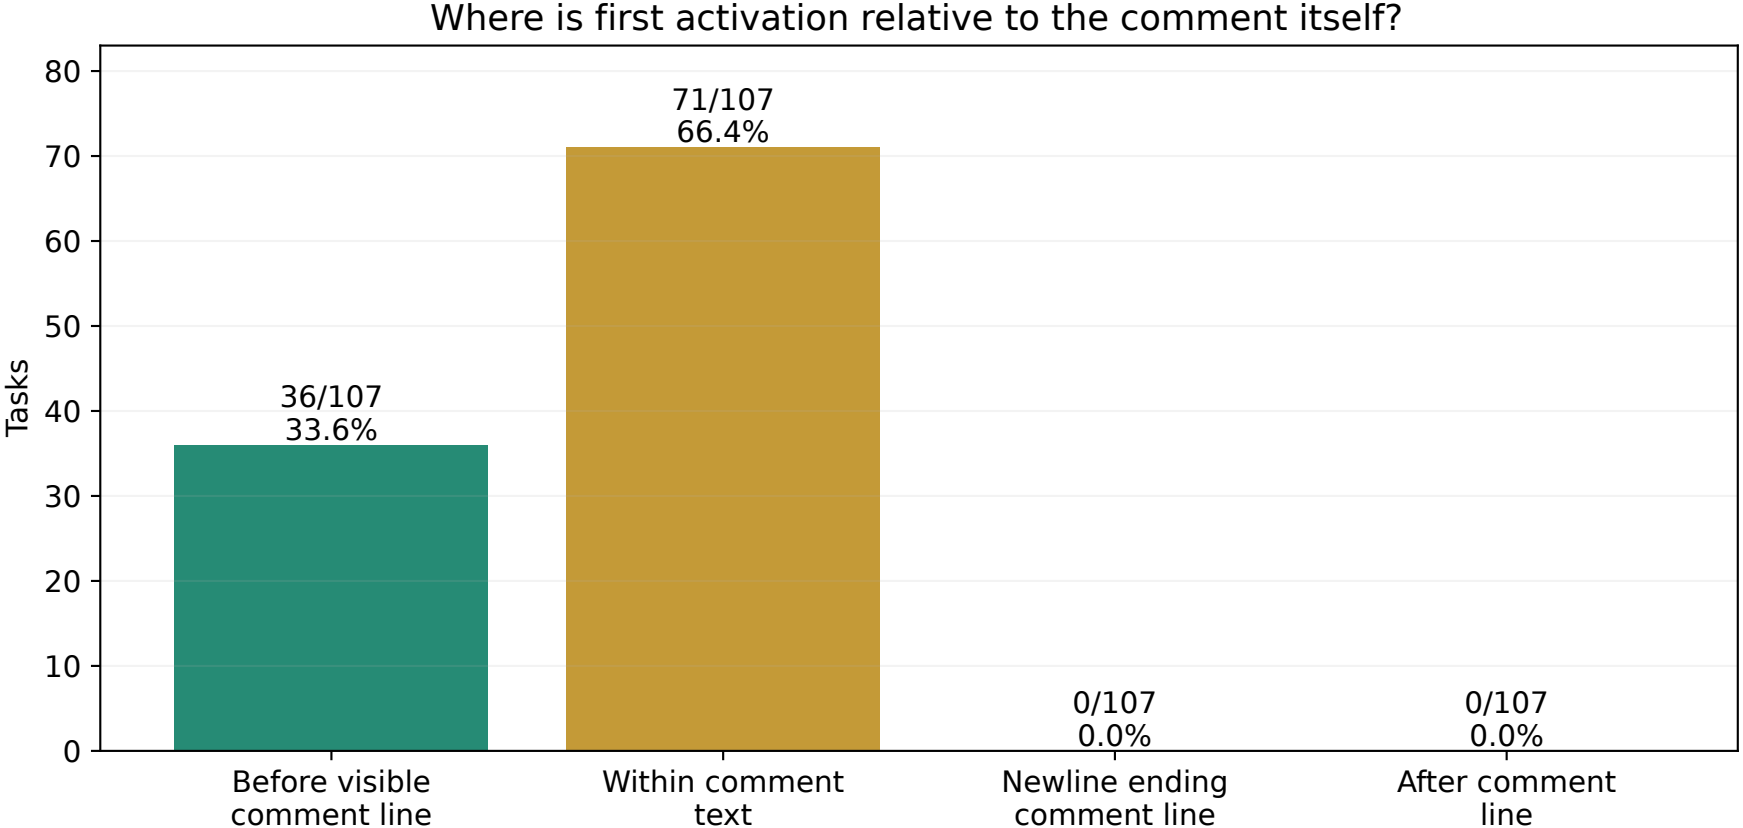

107 tasks whose first lexical marker is a test-file comment. Other 12 begin with task_func import / Test-class marker.
Comment text is not executable Python; this chart does not identify the first actual error.

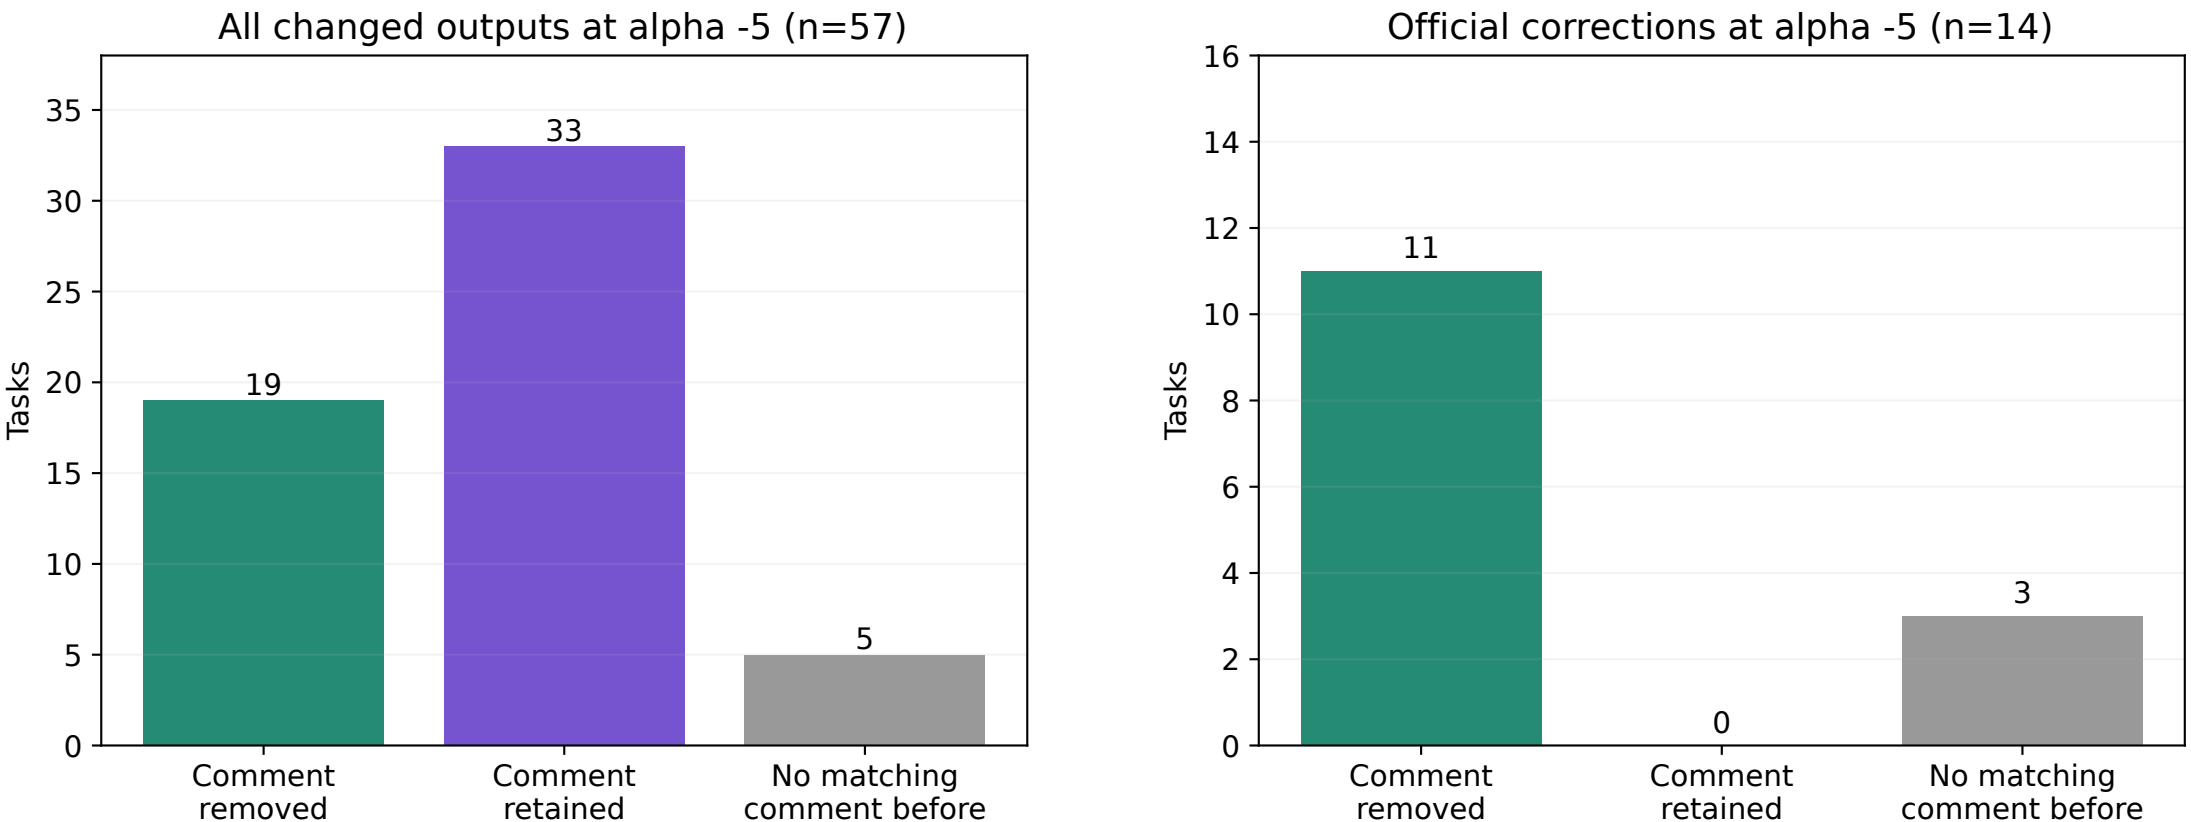

Matched alpha-zero versus alpha-minus-five evaluated code, same extractor. Removing a comment is not itself a repair.
All doses and task_func-import transitions are recorded separately in dose_summary.json and the case inspection files.

Removing the harmful test/import continuation is not the same as stopping generation

BigCodeBench/695 | baseline FAIL

```
#tests/test_task_func.py
import unittest
import numpy as np
from task_func import task_func

class TestTaskFunc(unittest.TestCase):
    def test_task_func(self):
        # Test case 1
```

main AST changed: False | display excerpt only

BigCodeBench/695 | steering PASS

```
#task_func([(1, 2, 3, 4), (5, 6, 7, 8), (9, 10, 11, 12)], 2)
```

main AST changed: False | display excerpt only

BigCodeBench/666 | baseline FAIL

```
#test_task_func.py
import unittest
from task_func import task_func

class TestTaskFunc(unittest.TestCase):
    def test_task_func(self):
        # Test case 1: Simple case
        seq = 'abc'
```

main AST changed: False | display excerpt only

BigCodeBench/666 | steering PASS

```
# Test the function
print(task_func('abc', {'a': 1, 'b': 2, 'c': 3}))
print(task_func('aabc', {'a': 10, 'b': -5, 'c': 3}))
```

main AST changed: False | display excerpt only

BigCodeBench/1115 | baseline FAIL

```
#test_task_func.py
import unittest
import random
from task_func import task_func

class TestTaskFunc(unittest.TestCase):
    def test_task_func(self):
        random.seed(0)
```

main AST changed: False | display excerpt only

BigCodeBench/1115 | steering PASS

```
# Write your code here

# Write your code here

# Write your code here

# Write your code here
```

main AST changed: False | display excerpt only

BigCodeBench/25 | baseline FAIL

```
#test_task.py
import unittest
import task

class TestTask(unittest.TestCase):
    def test_task_func(self):
        data_dict = {'key1': 'value1', 'key2': 'value2'}
        encoded_data = task.task_func(data_dict)
```

main AST changed: False | display excerpt only

BigCodeBench/25 | steering FAIL

```
#task_2.py
import base64
import json
import zlib

def task_func(encoded_data):
    """
    Decodes a base64 encoded string, decompresses it using zlib, and then
deserializes it to a dictionary.
```

main AST changed: False | display excerpt only

Exact saved evaluated-code excerpts. /695: cleanup; /666: executable examples remain; /1115: repetitive comments remain; /25: marker removal without a pass.
Full before/after code and diffs are retained in the companion inspection JSON / Markdown.

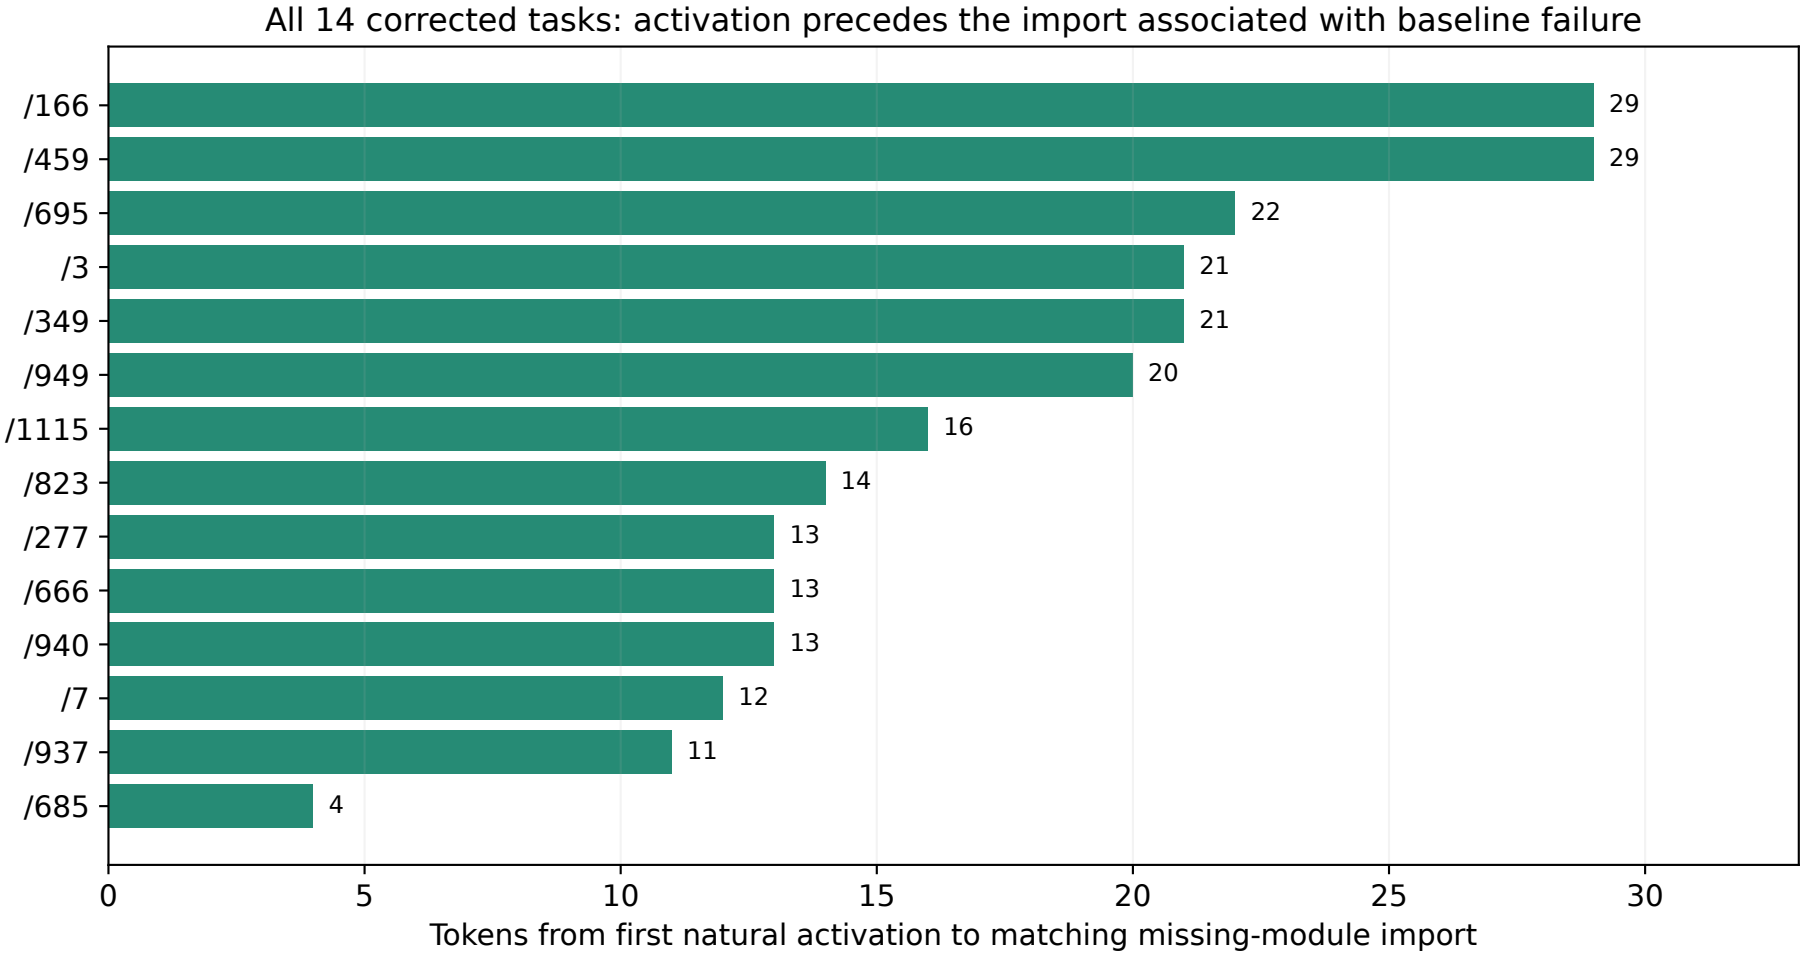

Baseline evaluator reports missing module task_func (11 cases) or task (3). Match earliest textual import of that module.
This is stronger temporal context than a comment marker, but not a proof of natural causal mediation. Frozen-text activation offsets.