DeepSeek 4810: first activation versus lexical contamination marker

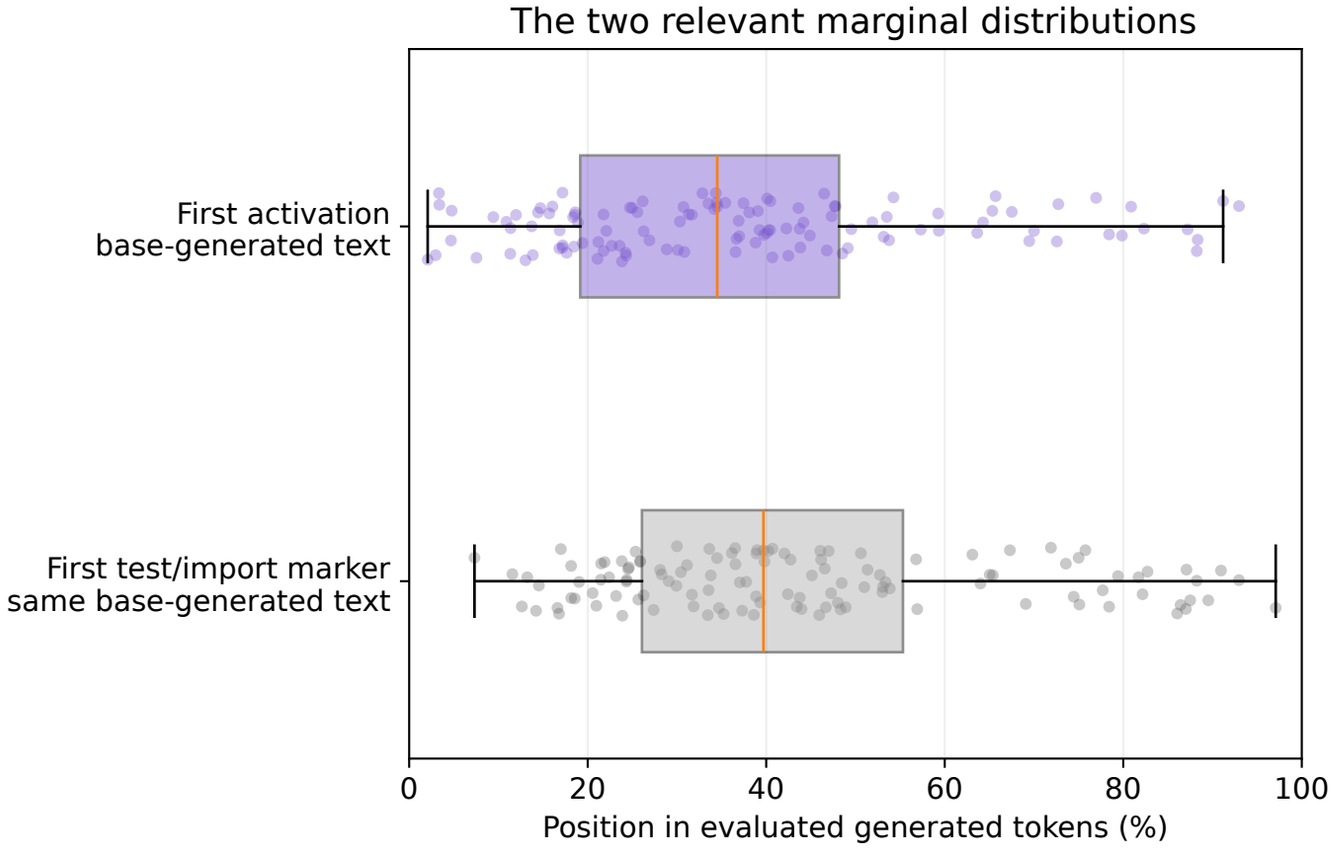

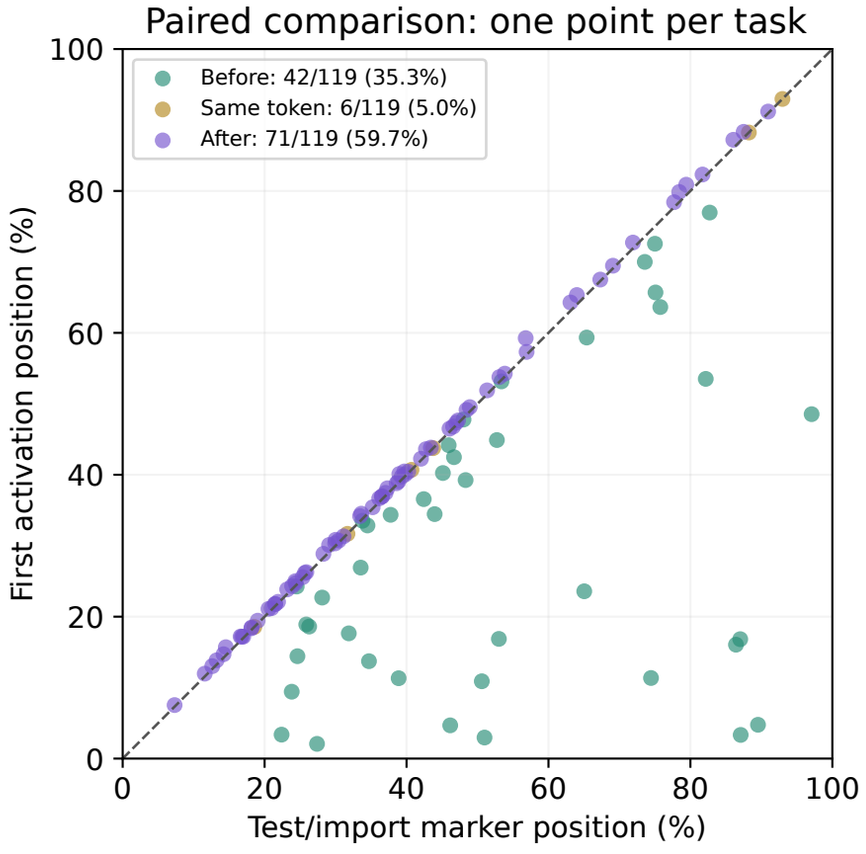

119 contamination-improvement tasks, original failing base text. Diagonal = same token position. Marker is not an annotated semantic error onset.

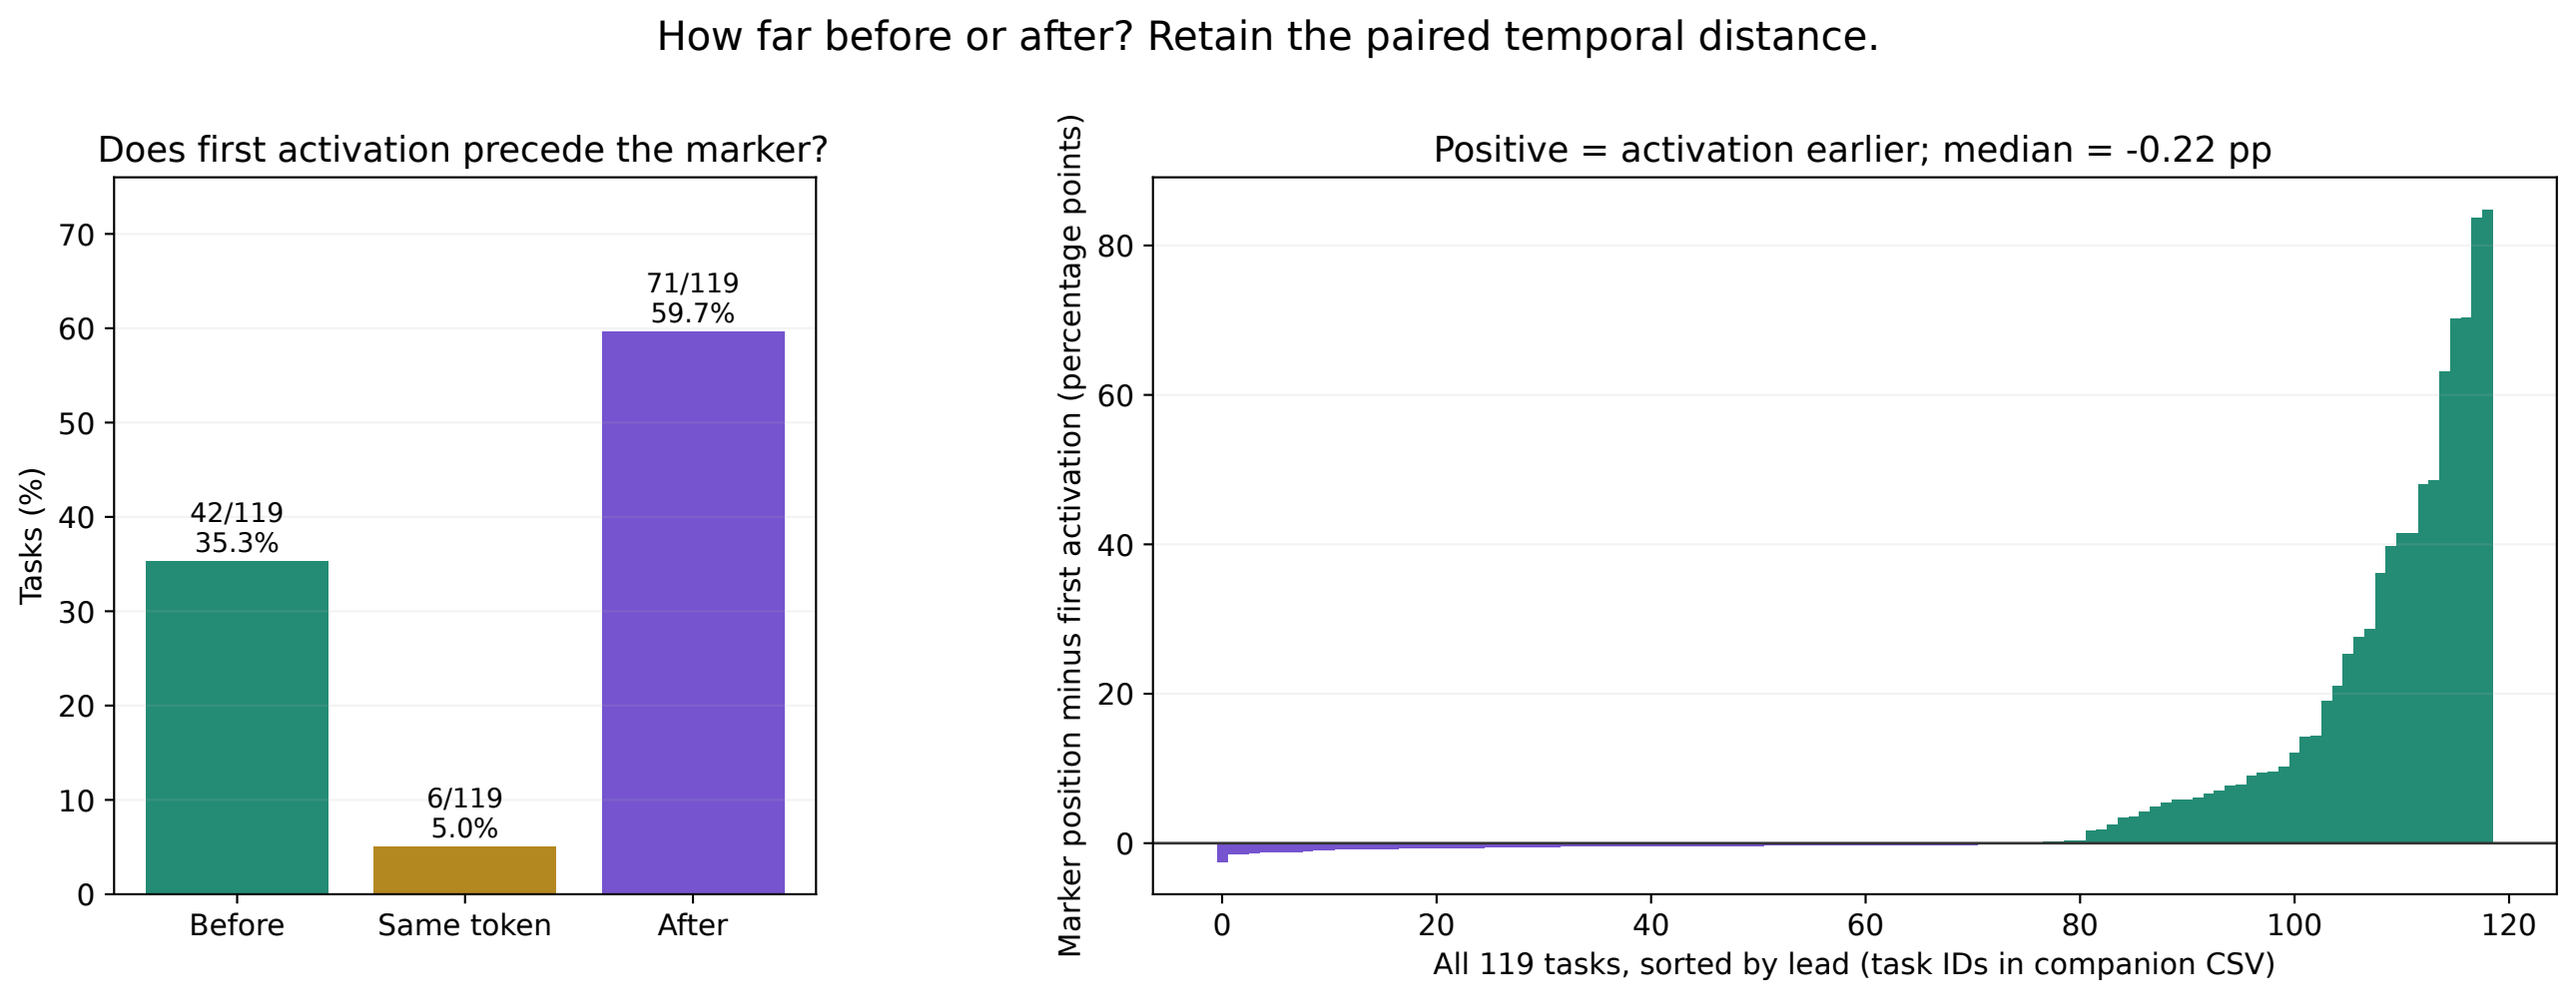

88/119 tasks within +/-5 percentage points; 99/119 within +/-10. Near the marker does not imply before it or causal mediation.

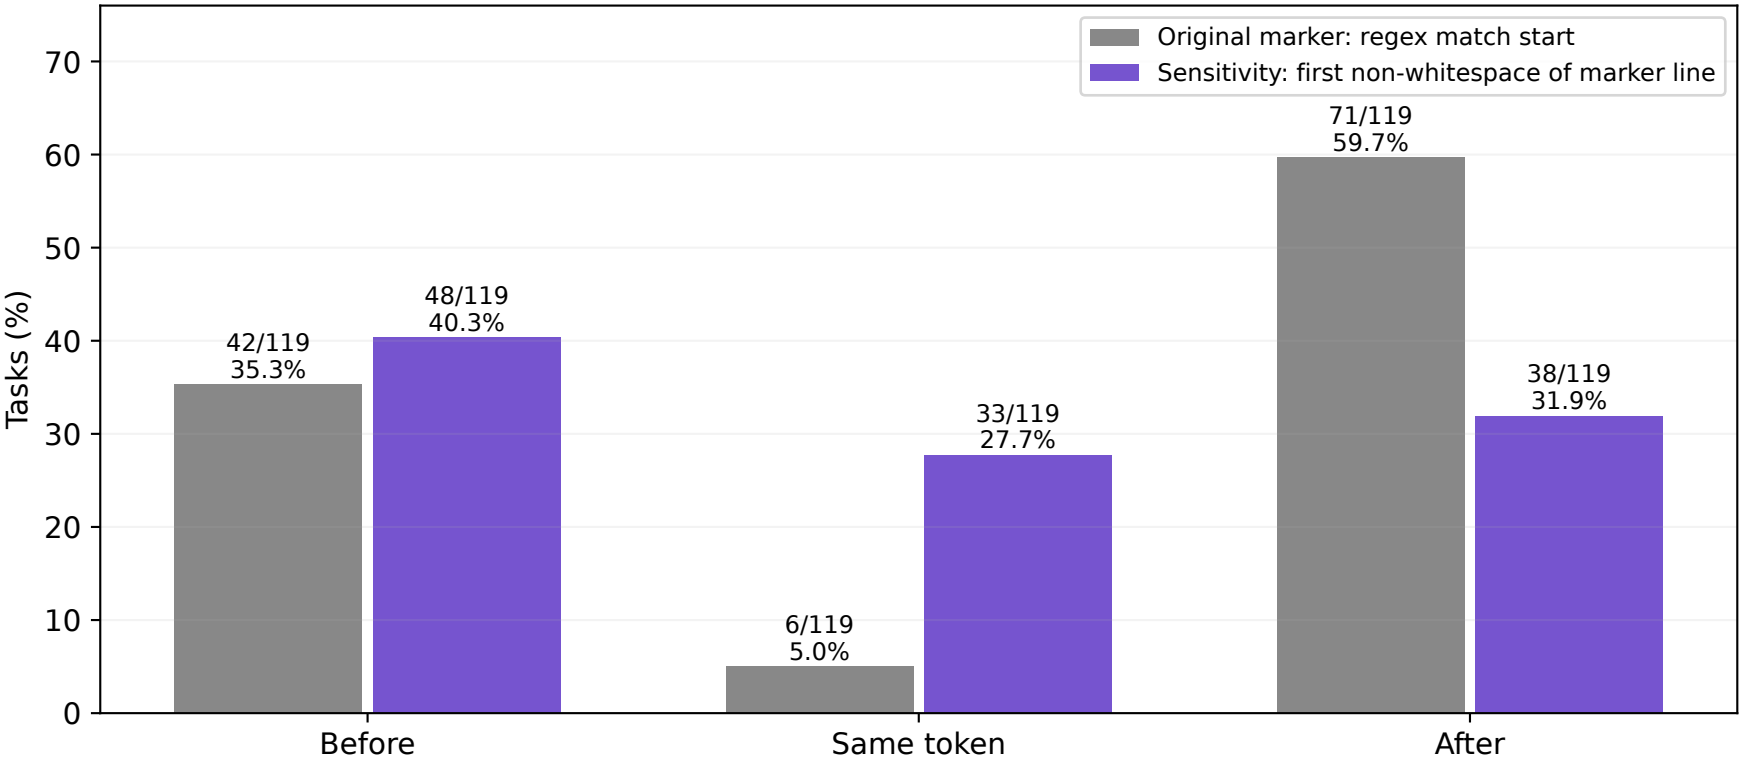

The original regex starts with ^\s*: it may include blank lines / indentation BEFORE the visible marker.
Sensitivity changes only the anchor, not the feature, cohort, model or code. Neither anchor is an annotated semantic error onset.