4810: does fine-tuning change the signal under identical text?

119 tasks x 2 source texts x 2 replayed models | layer 16 | existing activations only | CPU

One shared latent: z = TopK100(ReLU(u_base + u_FT + b)), with u_model = w_model dot h_model.
The two u terms below are contributions to one encoder, NOT independent sparse feature activations.

Original base-text region	Base contribution	FT contribution	FT lower tasks
Whole evaluated text	-0.524	-0.526	59/119
Marker +/-10 tokens	0.928	0.631	103/119
10 tokens before visible marker	-0.584	-0.570	55/119
Marker + next 10 tokens	2.303	1.723	110/119

Near the marker: own-branch decoder cosine is lower in FT in 113/119 tasks.
Same fixed decoder direction: lower in 119/119 (base reference) or 111/119 (FT reference).
The local model contrast is also more negative than the contrast outside the window.

Limits: selected exploratory cohort; geometric readouts are not calibrated latent activations.
A smaller representation under forced text does not prove that this feature mediates the behavioral improvement.
Bands on profile figures are task interquartile ranges, not confidence intervals; prompt activations excluded.