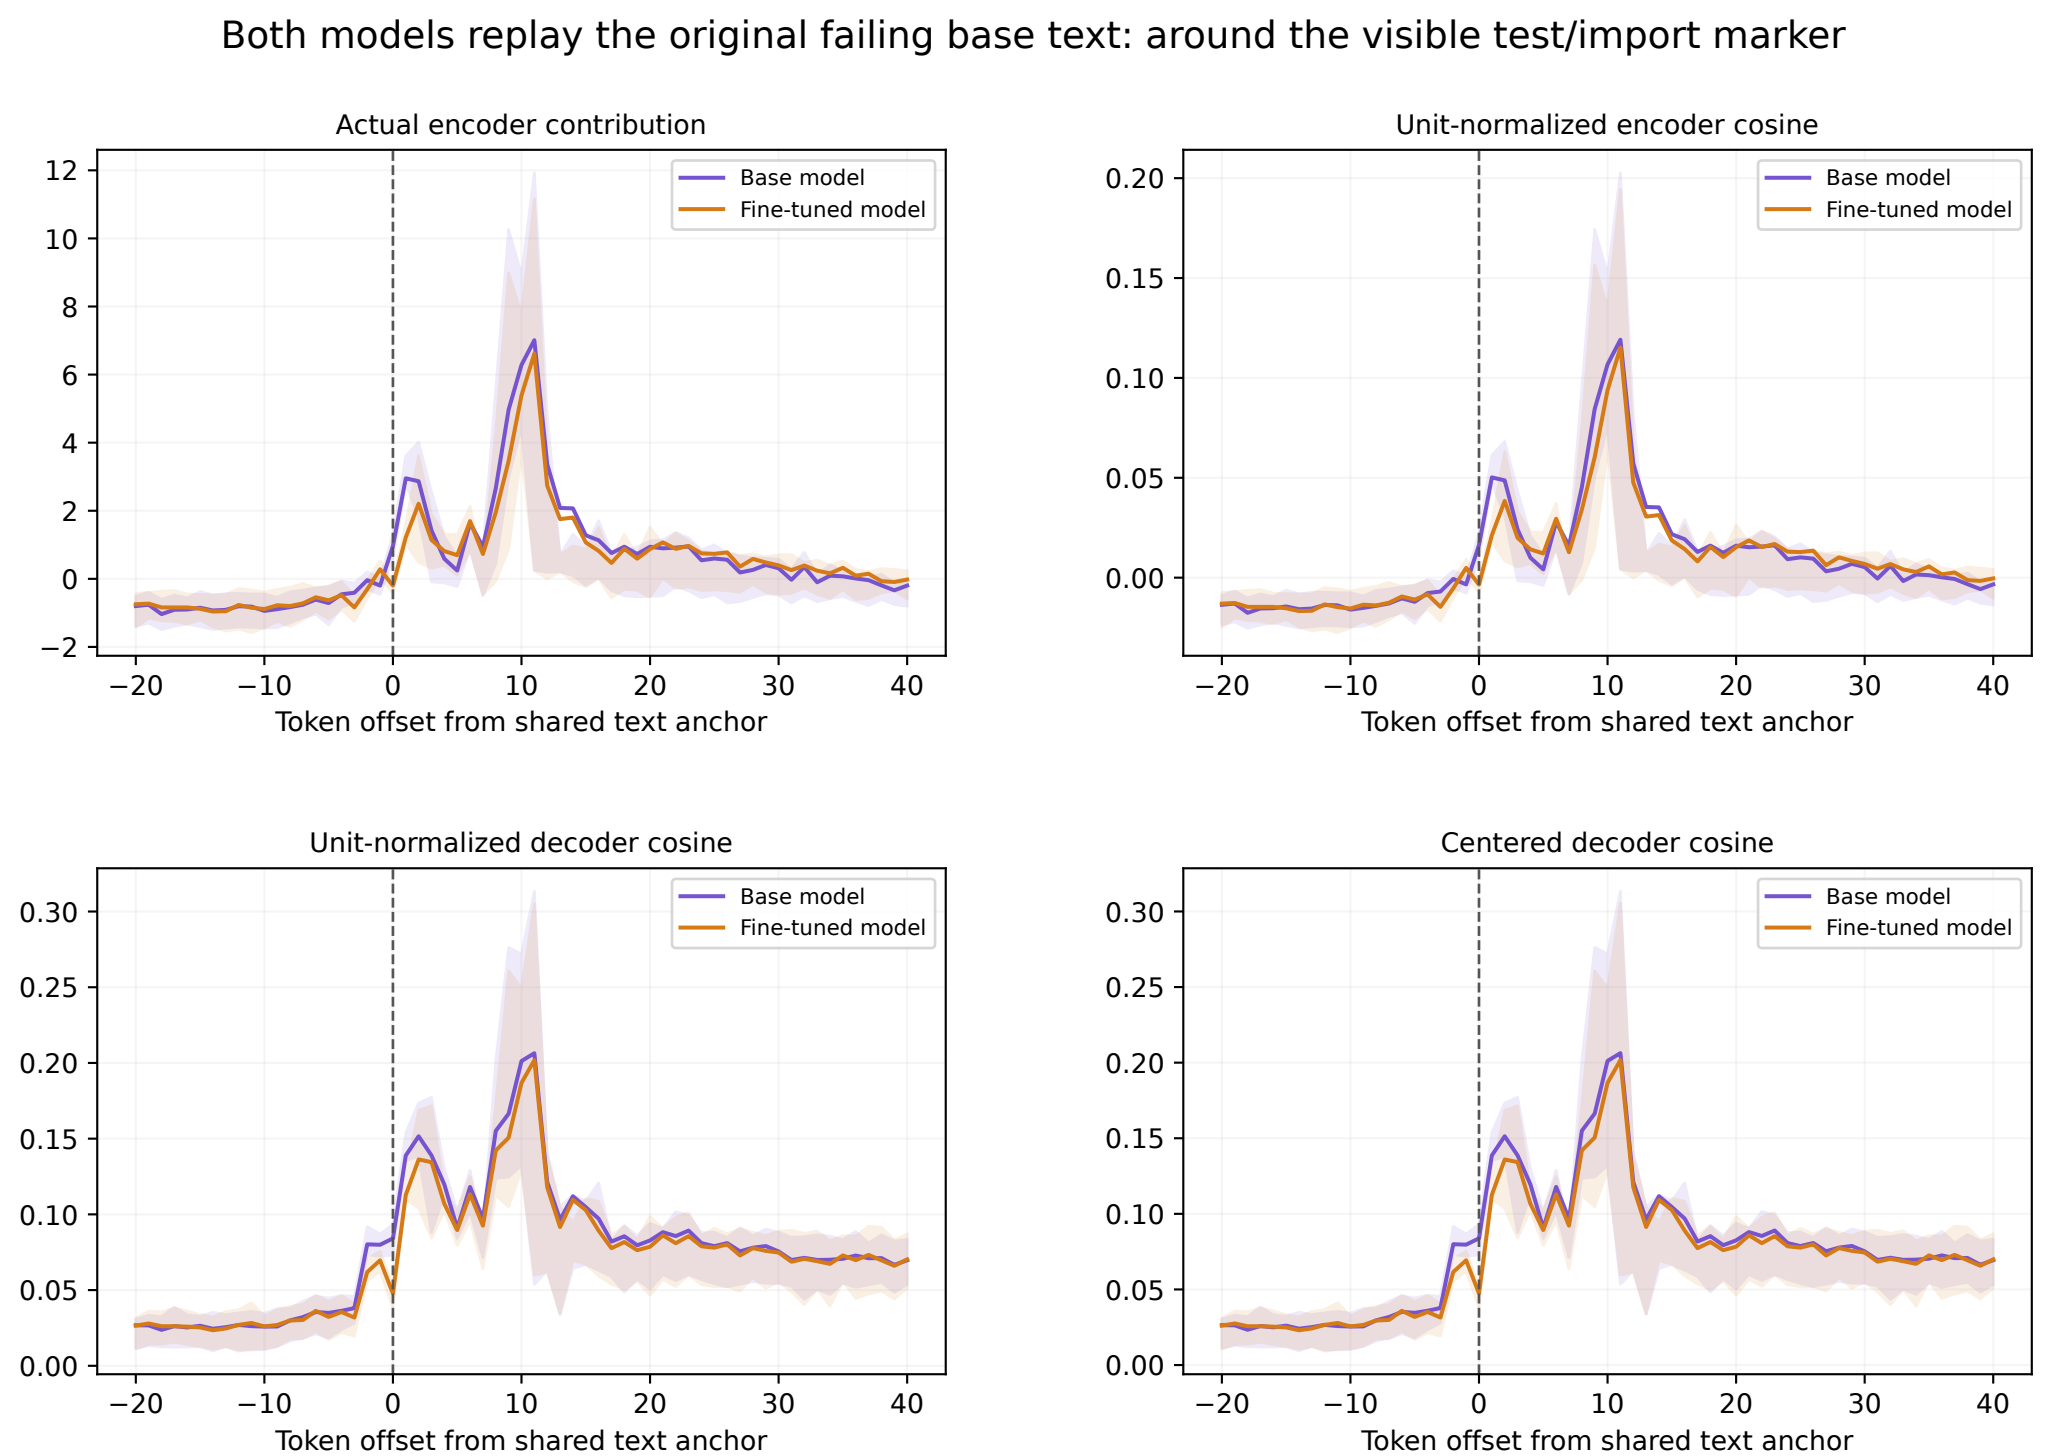
Both models replay the original failing base text: around the visible test/import marker
Actual encoder contribution
Unit-normalized encoder cosine
Unit-normalized decoder cosine
Centered decoder cosine
Base model
Fine-tuned model
Token offset from shared text anchor
Line: task-balanced mean; shaded band: interquartile task range, NOT confidence interval. Availability varies at sequence boundaries.
Marker is a lexical precursor, not necessarily the first harmful operation. Encoder terms sum with one shared bias before ReLU / joint TopK.