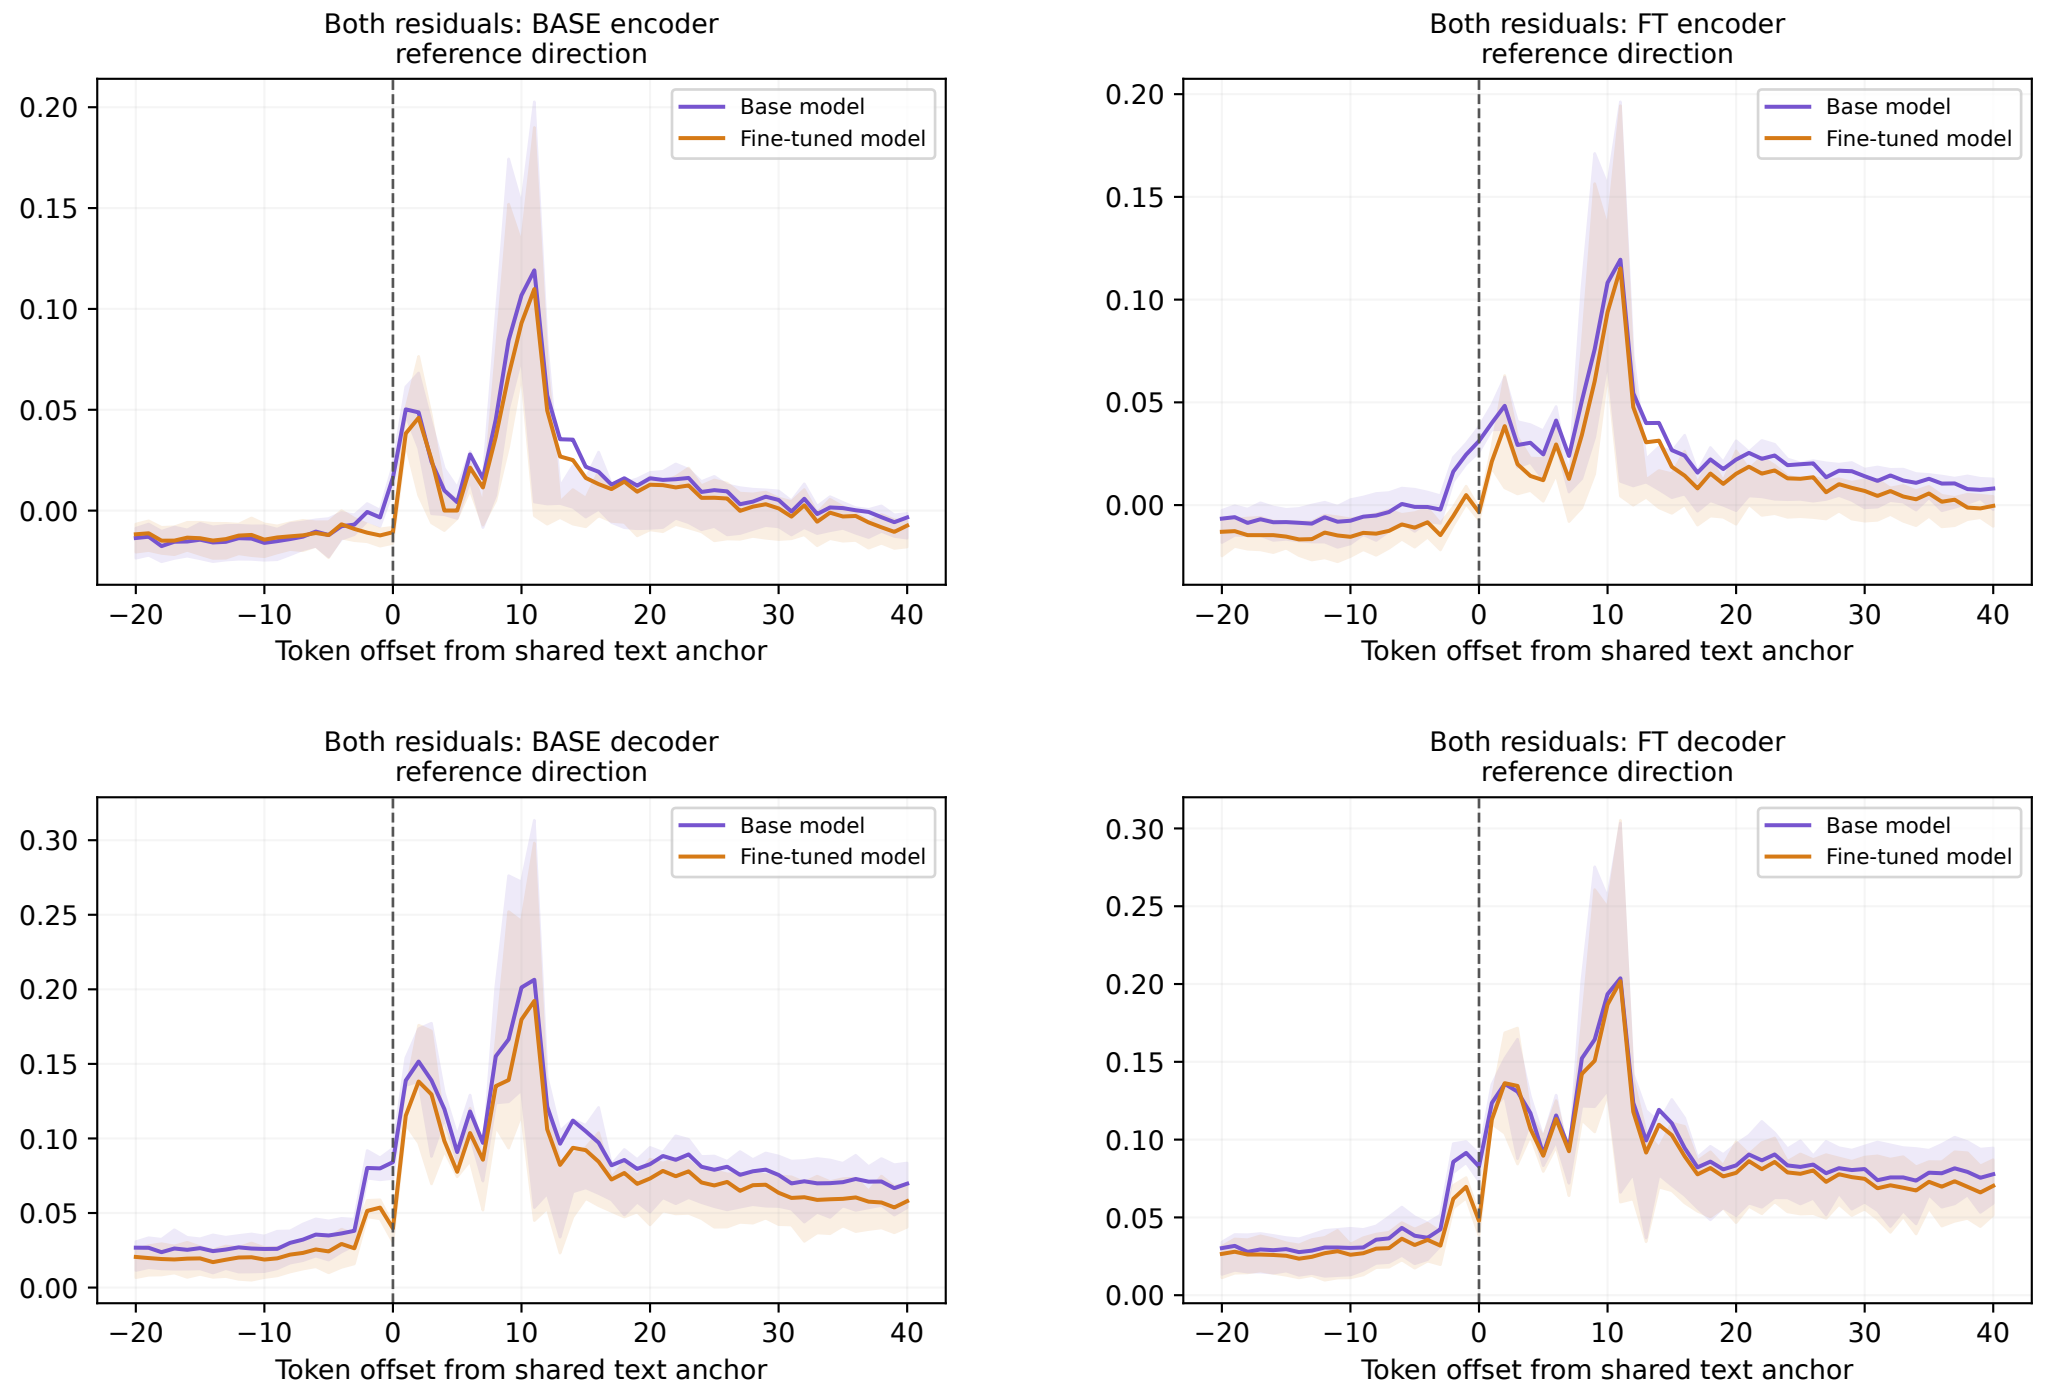
Sensitivity: same text AND the exact same readout vector for both models
Both residuals: BASE encoder reference direction
Both residuals: FT encoder reference direction
Both residuals: BASE decoder reference direction
Both residuals: FT decoder reference direction
Base model
Fine-tuned model
Token offset from shared text anchor
Original failing base texts, 119 paired tasks. Eliminates branch readout differences within each panel.
Same architecture / coordinate dimension; still a descriptive geometric comparison, not an independent sparse activation or causal mediation test.